

# IBM Assignment 3

October 11, 2022

```
[1]: import pandas as pd
import numpy as np
from sklearn.preprocessing import scale
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import f1_score
import seaborn as sns
import math
```

Load dataset

```
[2]: df = pd.read_csv("abalone.csv")
df = pd.DataFrame(df)
df.head()
```

```
[2]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	\
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	

	Shell weight	Rings
0	0.150	15
1	0.070	7
2	0.210	9
3	0.155	10
4	0.055	7

```
[3]: # created age column
df["Age"] = df["Rings"]+1.5
# dropped rings since unnecessary
df = df.drop("Rings", axis=1)
df
```

```
[3]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	\
0	M	0.455	0.365	0.095	0.5140	0.2245	
1	M	0.350	0.265	0.090	0.2255	0.0995	
2	F	0.530	0.420	0.135	0.6770	0.2565	
3	M	0.440	0.365	0.125	0.5160	0.2155	
4	I	0.330	0.255	0.080	0.2050	0.0895	
...	..	...	...	...	...	...	
4172	F	0.565	0.450	0.165	0.8870	0.3700	
4173	M	0.590	0.440	0.135	0.9660	0.4390	
4174	M	0.600	0.475	0.205	1.1760	0.5255	
4175	F	0.625	0.485	0.150	1.0945	0.5310	
4176	M	0.710	0.555	0.195	1.9485	0.9455	


	Viscera weight	Shell weight	Age
0	0.1010	0.1500	16.5
1	0.0485	0.0700	8.5
2	0.1415	0.2100	10.5
3	0.1140	0.1550	11.5
4	0.0395	0.0550	8.5
...	...	...	...
4172	0.2390	0.2490	12.5
4173	0.2145	0.2605	11.5
4174	0.2875	0.3080	10.5
4175	0.2610	0.2960	11.5
4176	0.3765	0.4950	13.5

[4177 rows x 9 columns]

Univariate analysis


```
[4]: age = df['Age'].value_counts().reset_index()
# barplot
sns.barplot(data=age, x='index', y='Age')
```

```
[4]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffaa3025450>
```




```
[5]: #histplot
sns.histplot(x=df.Age)
```

```
[5]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffaa2ec7150>
```



```
[6]: # boxplot
sns.boxplot(x=df.Age)
```


```
[6]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffaa29fd110>
```



Bivariate analysis


```
[7]: #barplot
sns.barplot(x=df['Diameter'], y=df.Age)
```

```
[7]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffaa2806b10>
```




```
[8]: #pointplot
sns.pointplot(x=df.Age, y=df.Height)
```

```
[8]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffaa289c3d0>
```



```
[9]: #scatter plot
sns.scatterplot(data=df,x="Age",y="Height")
```


```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x7ffaa29b3f10>
```



Multivariate analysis

```
[10]: #pairplot
sns.pairplot(data = df[["Height","Length","Diameter","Age","Whole_
↪weight","Shucked weight","Viscera weight","Shell weight"]])
```

```
[10]: <seaborn.axisgrid.PairGrid at 0x7ffaa227b150>
```



Descriptive statistics

```
[11]: df.describe()
```

```
[11]:
```

	Length	Diameter	Height	Whole weight	Shucked weight \
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367
std	0.120093	0.099240	0.041827	0.490389	0.221963
min	0.075000	0.055000	0.000000	0.002000	0.001000
25%	0.450000	0.350000	0.115000	0.441500	0.186000
50%	0.545000	0.425000	0.140000	0.799500	0.336000
75%	0.615000	0.480000	0.165000	1.153000	0.502000
max	0.815000	0.650000	1.130000	2.825500	1.488000

	Viscera weight	Shell weight	Age
count	4177.000000	4177.000000	4177.000000
mean	0.180594	0.238831	11.433684
std	0.109614	0.139203	3.224169
min	0.000500	0.001500	2.500000
25%	0.093500	0.130000	9.500000
50%	0.171000	0.234000	10.500000
75%	0.253000	0.329000	12.500000
max	0.760000	1.005000	30.500000

Missing values and how to deal with them

```
[12]: df.isnull().sum()
```

```
[12]: Sex                0
      Length            0
      Diameter          0
      Height            0
      Whole weight      0
      Shucked weight    0
      Viscera weight    0
      Shell weight      0
      Age               0
      dtype: int64
```

```
[13]: df.isna().sum()
      # no missing values
```

```
[13]: Sex                0
      Length            0
      Diameter          0
      Height            0
      Whole weight      0
      Shucked weight    0
      Viscera weight    0
      Shell weight      0
      Age               0
      dtype: int64
```

Find the outliers and replace them outliers

```
[14]: # replacing numerical outliers with lower and upper limits respectively

for i in df:
    if df[i].dtype=='int64' or df[i].dtypes=='float64':
        q1=df[i].quantile(0.25)
        q3=df[i].quantile(0.75)
```



```

iqr=q3-q1
upper=q3+1.5*iqr
lower=q1-1.5*iqr
df[i]=np.where(df[i] >upper, upper, df[i])
df[i]=np.where(df[i] <lower, lower, df[i])

```

Check for categorical columns and perform encoding

```

[15]: # identified and encoded the categorical values
from sklearn.preprocessing import LabelEncoder
encoder=LabelEncoder()
for i in df:
    if df[i].dtype=='object' or df[i].dtype=='category':
        print(i)
        df[i]=encoder.fit_transform(df[i])
df.head()

```

Sex

```

[15]:
  Sex  Length  Diameter  Height  Whole weight  Shucked weight \
0    2    0.455    0.365    0.095         0.5140         0.2245
1    2    0.350    0.265    0.090         0.2255         0.0995
2    0    0.530    0.420    0.135         0.6770         0.2565
3    2    0.440    0.365    0.125         0.5160         0.2155
4    1    0.330    0.255    0.080         0.2050         0.0895

  Viscera weight  Shell weight  Age
0          0.1010          0.150  16.5
1          0.0485          0.070   8.5
2          0.1415          0.210  10.5
3          0.1140          0.155  11.5
4          0.0395          0.055   8.5

```

Split the data into dependent and independent variables.

```

[16]: # independent variables
X = df.iloc[:, :-1].values
X

```

```

[16]: array([[2.      , 0.455 , 0.365 , ..., 0.2245, 0.101 , 0.15  ],
             [2.      , 0.35  , 0.265 , ..., 0.0995, 0.0485, 0.07  ],
             [0.      , 0.53  , 0.42  , ..., 0.2565, 0.1415, 0.21  ],
             ...,
             [2.      , 0.6    , 0.475 , ..., 0.5255, 0.2875, 0.308 ],
             [0.      , 0.625 , 0.485 , ..., 0.531 , 0.261 , 0.296 ],
             [2.      , 0.71  , 0.555 , ..., 0.9455, 0.3765, 0.495 ]])

```

```
[17]: # dependent variables
Y = df.iloc[:, -1].values
Y
df.head()
```

```
[17]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	\
0	2	0.455	0.365	0.095	0.5140	0.2245	
1	2	0.350	0.265	0.090	0.2255	0.0995	
2	0	0.530	0.420	0.135	0.6770	0.2565	
3	2	0.440	0.365	0.125	0.5160	0.2155	
4	1	0.330	0.255	0.080	0.2050	0.0895	

	Viscera weight	Shell weight	Age
0	0.1010	0.150	16.5
1	0.0485	0.070	8.5
2	0.1415	0.210	10.5
3	0.1140	0.155	11.5
4	0.0395	0.055	8.5

Scale independent variables

```
[18]: x = scale(df[["Viscera weight", "Length", "Diameter", "Height", "Whole weight",
↪ "Shucked weight", "Shell weight"]])
x
```

```
[18]: array([[ -0.73030425, -0.58311728, -0.44088378, ..., -0.6447403 ,
-0.61498531, -0.64518445],
[ -1.21388983, -1.46569411, -1.45976205, ..., -1.23820752,
-1.1916374 , -1.23138964],
[ -0.35725252,  0.04729474,  0.11949927, ..., -0.30943646,
-0.46736237, -0.20553056],
...,
[  0.98757595,  0.63567929,  0.67988232, ...,  0.71704585,
  0.77359293,  0.51257079],
[  0.74348037,  0.84581663,  0.78177015, ...,  0.54939393,
  0.79896563,  0.42464001],
[  1.80736865,  1.56028358,  1.49498494, ...,  2.30613921,
  2.71114397,  1.88282541]])
```

Split the data into training and testing

```
[19]: X = df.iloc[:, 1:7]
X
```

```
[19]:
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
0	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	0.350	0.265	0.090	0.2255	0.0995	0.0485
2	0.530	0.420	0.135	0.6770	0.2565	0.1415

3	0.440	0.365	0.125	0.5160	0.2155	0.1140
4	0.330	0.255	0.080	0.2050	0.0895	0.0395
...	...	...	...	...	...	...
4172	0.565	0.450	0.165	0.8870	0.3700	0.2390
4173	0.590	0.440	0.135	0.9660	0.4390	0.2145
4174	0.600	0.475	0.205	1.1760	0.5255	0.2875
4175	0.625	0.485	0.150	1.0945	0.5310	0.2610
4176	0.710	0.555	0.195	1.9485	0.9455	0.3765

[4177 rows x 6 columns]

```
[20]: Y = df.iloc[:, -1]
      Y
```

```
[20]: 0      16.5
      1       8.5
      2     10.5
      3     11.5
      4       8.5
      ...
      4172    12.5
      4173    11.5
      4174    10.5
      4175    11.5
      4176    13.5
      Name: Age, Length: 4177, dtype: float64
```

```
[21]: # splitting to train and test data
      x_train,x_test,y_train,y_test = train_test_split(X, Y, test_size=0.
      ↪25,random_state =42)
```

Build the Model

```
[22]: # linear regression model
      model = LinearRegression()
```

Train the model

```
[23]: model.fit(x_train,y_train)
```

```
[23]: LinearRegression()
```

Test the model

```
[24]: y_predict = model.predict(x_test)
      y_predict
```

```
[24]: array([12.80459703, 11.53324712, 15.38929967, ..., 13.49727819,  
          11.82966664, 10.72212477])
```

```
[25]: print("Mean Squared Error: ",mean_squared_error(y_test, y_predict))  
      print("Root Mean Squared Error: ",math.sqrt(mean_squared_error(y_test, y_↵  
      ↵y_predict)))
```

Mean Squared Error: 3.4158644158621536

Root Mean Squared Error: 1.848205728771057